

COMP 671D: Homework 3

Author: Samuel Eure

Problem 3

```
In [1]: import numpy as np
```

February 17, 2019

```
import pandas as pd
import matplotlib.pyplot as plt
import scipy as scipy
from scipy.optimize import minimize
import random
from sklearn import metrics
```

```
In [2]: train      = pd.read_csv("house_votes_84_train.csv", sep = ',')
test      = pd.read_csv("house_votes_84_test.csv", sep = ',')
samples   = len(train.iloc[:,0])
```

```
In [3]: def getJandError(train,d):
    j_t = -1
    minError = 2.0
    realY = np.array(train.iloc[:,0])
    for j in range(len(train.iloc[0,1:])):
        j += 1
        predictions = np.array(train.iloc[:,j])
        score = np.matmul((predictions != realY),d)
        if score < minError:
            minError = score
            j_t = j
    return(j_t, minError)
```

Problem 3a

AdaBoost Implementation

```
In [4]: T = len(train.iloc[0,1:])
realY = np.array(train.iloc[:,0])
d = np.ones(samples)/samples
weights = np.zeros(T)
```

```

for t in range(100):    # I run the algorithm 100 times however my weights
                        # essentially converge after the first iteration

    j_t, error = getJandError(train,d)

    alpha = 0.5*np.log((1-error)/error)
    weights[j_t-1] += alpha
    predictions = np.array(train.iloc[:,j_t])
    correct     = (predictions == realY)*np.exp(-1*alpha)
    incorrect   = (predictions != realY)*np.exp(alpha)
    d = correct + incorrect
    d = d/sum(d)

    features = np.array(train.iloc[:, 1:])

```

Adaboost Model Evaluation

```

In [5]: print("Training Accuracy:",np.mean(np.sign(np.matmul(
        np.array(train.iloc[:,1:]),weights))==realY)) #Accuracy
        print("Test Accuracy:",np.mean(np.sign(np.matmul(
        np.array(test.iloc[:,1:]),weights))==np.array(test.iloc[:,0]))) #Accuracy

```

Training Accuracy: 0.9428571428571428
Test Accuracy: 0.98

Problem 3b

Logistic Regression Implementation

For this problem, I will find the minimum of the negative log likelihood instead of the maximum log likelihood. This will accomplish the same thing.

Given that: $-\log(P(Y_i = y_i | \vec{\beta}, \vec{x}_i)) = \log(1 + e^{-y_i \vec{\beta} \cdot \vec{x}_i})$,

the negative log likelihood of $-\log(L(\beta))$ is equal to $-\log(L(\beta)) = \sum_{i=1}^n \log(1 + e^{-y_i \vec{\beta} \cdot \vec{x}_i})$.

This function will be minimized with respect to β

```

In [6]: def getNLL(betas):
        data = np.array(train)
        prod = np.matmul(data[:, 1:], betas) #features
        yValues = np.array(data[:, 0]) #real y values
        yProd   = yValues*prod
        return(np.sum(np.log(1+np.exp( -1*yProd))))

        betas0 = np.zeros(T)
        getNLL(betas0)    #Algorithm Testing

```

Out [6]: 266.8616645155789

Determining "optimal" beta values

```
In [7]: results = scipy.optimize.minimize(getNLL, x0 = np.zeros(T))
        optimalBetas = results.x
        print("Optimal Beta Values:", optimalBetas)
```

```
Optimal Beta Values: [ 0.07097047 -0.77902062 -1.34247057  4.29392102  1.78724924 -2.19311427
 1.7240725  -0.75573081 -1.05317142  1.4752639  -1.43506075  1.08354749
-0.32326115 -0.76060663 -0.29633747 -0.83586897]
```

Model Evaluation for Logistic Regression

```
In [8]: def evaluateBetas(data, optimalBetas):
        data = np.array(data)
        y = data[:, 0]
        X = data[:, 1:]
        BX = np.matmul(X, optimalBetas)

        probYis1 = 1/(1+ np.exp(-1*BX))
        prePreds = probYis1 >= 0.5
        predictions = prePreds*2 -1
        fpr,tpr, thresholds = metrics.roc_curve(y,predictions, pos_label=1)
        print("AUC:", np.round(metrics.auc(fpr, tpr),4))
        acc = np.mean((predictions == y))
        print("Accuracy:", 100*np.round(acc,4), "%")
        return()

        print("For the training data set:")
        evaluateBetas(train, optimalBetas)
        print("-----")
        print("For the test data set:")
        evaluateBetas(test, optimalBetas)
```

For the training data set:

AUC: 0.9677

Accuracy: 96.36 %

For the test data set:

AUC: 0.9228

Accuracy: 94.0 %

```
Out[8]: ()
```

Problem 3c

Model Comparison

In this problem, I will compare the test accuracies, ROC Curves/AUC Values, and F1 scores.

```
In [9]: test = np.array(test)
        y = test[:, 0]
        X = test[:, 1:]

        # Logistic Regression Data
        BX = np.matmul(X, optimalBetas)
        LR_probs = 1/(1+ np.exp(-1*BX))
        LR_predictions = (LR_probs >= 0.5)*2 -1

        # AdaBoost Data
        AD_rawScores = np.matmul(np.array(X),weights) #Used for ROC curve
        AD_predictions = np.sign(AD_rawScores)

In [10]: def getROC(name, probabilities):
          fpr, tpr, thresh = metrics.roc_curve(y,probabilities , pos_label = 1)
          auc = np.round(metrics.auc(fpr, tpr),4)

          plt.plot(fpr, tpr)
          plt.title("ROC Curve for " + name + " (AUC: "+str(auc)+")", fontsize = 15)
          plt.plot([0,1],[0,1],
                   color = 'black',
                   alpha = .5,
                   linestyle = 'dashed')
          plt.xlabel("FPR", fontsize = 15)
          plt.ylabel("TPR", fontsize = 15)

          def getAcc_F1(predictions, y):
              TotPos = sum(y == 1)
              TotNeg = len(y) - TotPos

              correct = (predictions == y)
              PredPos = (predictions == 1)


              TP = sum(PredPos*correct)
              recall = TP/TotPos
              precision = TP/sum(PredPos)
              F1 = 2*((precision*recall)/(precision+recall))
              print("Accuracy: ",np.mean(correct))
              print('Recall: ',recall)
              print("Precision:", np.round(precision,4))
              print("F1 Score: ", np.round(F1,4))

          getAcc_F1(AD_predictions, y)

In [11]: #Logistic Regression
```


```
getROC('Logistic Reg.', LR_probs)
getAcc_F1(LR_predictions,y)
```

Accuracy: 0.94
Recall: 0.875
Precision: 0.9333
F1 Score: 0.9032



```
In [12]: #AdaBoost
getROC('AdaBoost',AD_rawScores)
getAcc_F1(AD_predictions,y)
```

Accuracy: 0.98
Recall: 1.0
Precision: 0.9412
F1 Score: 0.9697



Comments on the models

It appears that the Adaboost algorithm obtained a higher accuracy, recall, precision, and F1 score than the Logistic Regression algorithm. This suggests that the Adaboost algorithm does a better job in classification than the logistic regression algorithm for this data set and given the specifications of the problem. However, it should be noted that the Logistic Regression algorithm obtained a higher area under the ROC curve (i.e. AUC score) than the Adaboost algorithm, although only slightly (about a 0.0128 difference).

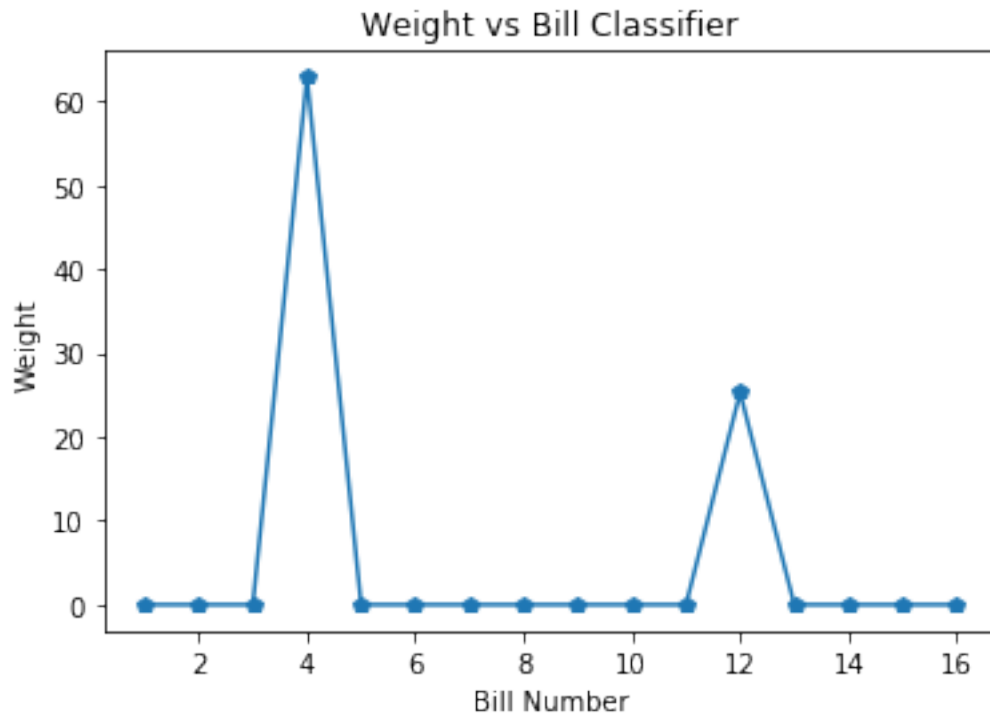
In summary, the Adaboost algorithm seemed to perform better than the logistic regression algorithm, however both models did decently well in the classification problem given that they both have accuracies in the mid to high nineties.

Problem 3d

During each iteration of the AdaBoost algorithm one of sixteen classifiers was chosen and coordinate gradient descent was performed in the direction of that classifiers. Given that the algorithm chooses between 16 weak classifiers at each step, the weight vector for the AdaBoost implementation for this data set is 16 dimensional, with the i^{th} dimension being attributed to the i^{th} weak classifier, which is related to the i^{th} bill. Thus, the AdaBoost algorithm would suggest that those classifiers with large weights are "important" for making predictions. Thus the most important bills for making predictions from the view of the Adaboost algorithm are...

```
In [13]: bills = [1,2,3,4,5,6,7,8,9,10,11,12,13,14,15,16]
plt.plot(bills,weights, '-p')
plt.title("Weight vs Bill Classifier")
plt.xlabel("Bill Number")
plt.ylabel("Weight")
```

```
Out[13]: Text(0, 0.5, 'Weight')
```



Bills 4 and 12 are the most significant predictors given that they're the only bills with nonzero weights attribute to them. Indeed, Adaboost has determined that given these two variable, the remaining 14 variables are essentially superfluous. Furthermore, it appears that even when the 12th bill is removed as a predictor, the Adaboost algorithm is able to produce the same accuracy as it was able to with all the predictors, so Bill 12 doesn't seem that important in terms of predictions in comparison to Bill 4. This is shown below:

```
In [14]: train = pd.DataFrame(train)
test = pd.DataFrame(test)
print("Original Training Accuracy:", np.mean(np.sign(np.matmul(
    np.array(test.iloc[:,1:]),weights))==np.array(test.iloc[:,0]))) #Accuracy
remember12 = np.array(train.iloc[:, 12])
train.iloc[:, 12] = 1 #rendering it useless
T = len(train.iloc[0,1:])
realY = np.array(train.iloc[:,0])
d = np.ones(samples)/samples
```

```

weights = np.zeros(T)
for t in range(100):    # I run the algorithm 100 times however my weights
                        # essentially converge after the first iteration

    j_t, error = getJandError(train,d)

    alpha = 0.5*np.log((1-error)/error)
    weights[j_t-1] += alpha
    predictions = np.array(train.iloc[:,j_t])
    correct    = (predictions == realY)*np.exp(-1*alpha)
    incorrect  = (predictions != realY)*np.exp(alpha)
    d = correct + incorrect
    d = d/sum(d)

print("Test Accuracy when 12 is removed:",np.mean(np.sign(np.matmul(
    np.array(test.iloc[:,1:]),weights))==np.array(test.iloc[:,0]))) #Accuracy

train.iloc[:, 12] = pd.DataFrame(remember12)
remember4  = np.array(train.iloc[:,4])
train.iloc[:,4] = 1
d = np.ones(samples)/samples
weights = np.zeros(T)
for t in range(100):    # I run the algorithm 100 times however my weights
                        # essentially converge after the first iteration

    j_t, error = getJandError(train,d)

    alpha = 0.5*np.log((1-error)/error)
    weights[j_t-1] += alpha
    predictions = np.array(train.iloc[:,j_t])
    correct    = (predictions == realY)*np.exp(-1*alpha)
    incorrect  = (predictions != realY)*np.exp(alpha)
    d = correct + incorrect
    d = d/sum(d)

print("Test Accuracy when 4 is removed:",np.mean(np.sign(np.matmul(
    np.array(test.iloc[:,1:]),weights))==np.array(test.iloc[:,0]))) #Accuracy
train.iloc[:,4] = pd.DataFrame(remember4)

```


Original Training Accuracy: 0.98
 Test Accuracy when 12 is removed: 0.98
 Test Accuracy when 4 is removed: 0.86

This suggests that Bill 4 is the most significant predictor of partisan ship from an Adaboost perspective, however, in it's absence, Adaboost is still able to obtain an 0.86 accuracy, suggesting that some of the remaining bills contain predictive power of their own.

From the standpoint of logistic regression, it would appear that bills with large corresponding beta values would be important predictors of partisanship. Let us look at the which beta values are largest:

```
In [15]: plt.plot(bills,optimalBetas, '-p')
plt.title("Beta Values vs Bill")
plt.xlabel("Bill Number")
plt.ylabel("Beta Value")
```

```
Out[15]: Text(0, 0.5, 'Beta Value')
```



If the magnitude of the beta value was attributed to the importance of the predictions, it would appear that bills number 4, 5, 7, 10, and 12 are the most important. To test this hypothesis, I will "remove" the information stored in each of these bills by setting their values equal to 1 for each observation, and then run my logistic regression algorithm on this new data set for each of the features:

```
In [16]: print("Original Data Set")
evaluateBetas(test, optimalBetas)
print("-----")
for b in bills:
    removedBill = np.array(train.iloc[:, b])
    train.iloc[:,b] = 1 #rendering the bill useless
    results = scipy.optimize.minimize(getNLL, x0 = np.zeros(T))
```

```

optimalBetas = results.x
print("Accuracy and AUC when bill",str(b), "is removed:")
evaluateBetas(test, optimalBetas)
train.iloc[:,b] = pd.DataFrame(removedBill)
print("-----")

```

Original Data Set

AUC: 0.9228

Accuracy: 94.0 %

Accuracy and AUC when bill 1 is removed:

AUC: 0.9081

Accuracy: 92.0 %

Accuracy and AUC when bill 2 is removed:

AUC: 0.8327

Accuracy: 84.0 %

Accuracy and AUC when bill 3 is removed:

AUC: 0.9393

Accuracy: 94.0 %

Accuracy and AUC when bill 4 is removed:

AUC: 0.7574

Accuracy: 76.0 %

Accuracy and AUC when bill 5 is removed:

AUC: 0.8934

Accuracy: 90.0 %

Accuracy and AUC when bill 6 is removed:

AUC: 0.8768

Accuracy: 90.0 %

Accuracy and AUC when bill 7 is removed:

AUC: 0.9853

Accuracy: 98.0 %

Accuracy and AUC when bill 8 is removed:

AUC: 0.9393

Accuracy: 94.0 %

Accuracy and AUC when bill 9 is removed:

AUC: 0.954

Accuracy: 96.0 %

Accuracy and AUC when bill 10 is removed:

AUC: 0.9853

Accuracy: 98.0 %

Accuracy and AUC when bill 11 is removed:

AUC: 0.9393

Accuracy: 94.0 %

Accuracy and AUC when bill 12 is removed:

AUC: 0.8787

Accuracy: 88.0 %

Accuracy and AUC when bill 13 is removed:

AUC: 0.8621

Accuracy: 88.0 %

Accuracy and AUC when bill 14 is removed:

AUC: 0.9062

Accuracy: 94.0 %

Accuracy and AUC when bill 15 is removed:

AUC: 0.9393

Accuracy: 94.0 %

Accuracy and AUC when bill 16 is removed:

AUC: 0.8768

Accuracy: 90.0 %

These results suggest that removing features 4 dramatically reduces the test accuracy and removing feature 2 someone significantly reduces the test accuracy. This would suggest that each of these bills are important for predictions in general. Another interesting discovery of these calculations is that the removal of bills 7,9, and 10 actually leads to an INCREASE in test accuracy. This is most likely due to overfitting of the training data when using all bills are predictors.

In []: